

```

+-----+
|               CS5600               |
|   PROJECT 3: VIRTUAL MEMORY   |
|             DESIGN DOCUMENT             |
+-----+

```

---- GROUP ----

>> Fill in the names and email addresses of your group members.

```

Wanting Jiang          <jiang.wa@husky.neu.edu>
<wanting>
Heng Xu                <xu.he@husky.neu.edu>
<hengx>
Douglas Johnston       <johnston.5douglas@gmail.com>
<douglasj>

```

---- PRELIMINARIES ----

>> If you have any preliminary comments on your submission,  
notes for the

>> TAs, or extra credit, please give them here.

>> Please cite any offline or online sources you consulted while  
>> preparing your submission, other than the Pintos  
documentation, course  
>> text, lecture notes, and course staff.

Refer to this online sources:

<https://github.com/ryantimwilson/Pintos-Project-3>

```

PAGE TABLE MANAGEMENT
=====

```

---- DATA STRUCTURES ----

>> A1: Copy here the declaration of each new or changed `struct`  
or  
>> `struct` member, global or static variable, `typedef`, or  
>> enumeration. Identify the purpose of each in 25 words or  
less.

In "page.h":

```

// spt_entry is the supplement page which records the
information

```

```

// of the user page
struct spt_entry {
    void *upage;

```

```

uint8_t *frame;          // frame entry
struct file *file;       // file
off_t ofs;               // file offset
bool writable;           // whether file can be writable or not
uint32_t read_bytes;     // read bytes
uint32_t zero_bytes;     // zero bytes
bool loaded;             // whether page is loaded or not
uint32_t swap_sector;    // if page is swapped, point out the
swap sector
bool swap;               // whether page is swapped or not
bool mmap;               // whether it is a memory mapped file or
not
int mapid;               // if it is a memory mapped file, point
out the map id
bool pinned;             // prevent other processes from
accessing it when page is being used

    struct hash_elem elem;
};

```

In "frame.h":

```

struct list frame_table;
struct lock frame_lock;

struct frame_entry {
    uint8_t *frame;          // frame address
    struct spt_entry *spte;  // page entry
    struct thread *owner;    // thread id
    struct list_elem elem;
};

```

In "thread.h", we added the following fields:

```

struct hash spt; /* Supplemental page table. */
struct list mmap_list; //a list of mmap_entry
int mapid; // map id, used when creates a new mmap_entry

```

---- ALGORITHMS ----

```

>> A2: In a few paragraphs, describe your code for locating the
frame,
>> if any, that contains the data of a given page.

```

There are 4 steps to locate the frame:

1. Round down the user virtual address to the corresponding memory page address.
2. Then call `palloc_get_frame()` to allocate a frame for this supplemental page. If no free frame is available right now, call `frame_evict()` to free a frame.
3. Call `install_page()` to map user virtual address to physical frame address.
4. Read data from the file if user virtual address is from a file or `mmap`, or from the swap partition if user virtual address is from swap.

>> A3: How does your code coordinate accessed and dirty bits between  
>> kernel and user virtual addresses that alias a single frame,  
or  
>> alternatively how do you avoid the issue?  
We start by checking the pointer to make sure it is valid in `syscall.c` and `exception.c`. We record the user virtual address in "frame" and "upage" to find the supplemental page table, and kernel virtual address is never used.

---- SYNCHRONIZATION ----

>> A4: When two user processes both need a new frame at the same time,  
>> how are races avoided?

We have a new struct `frame_lock` and added it to three functions:

1. `add_to_frame_table();`
2. `free_frame();`
3. `frame_evict();`

The lock ensures that at any given time, only one process can acquire the lock and get page index. So when two processes successfully obtain pages, they will add the new frame to the frame table sequentially since we add lock in `add_to_frame_table()`. Same condition happens to `free_frame()` and `frame_evict()`. When two or more processes get pages, every action on frame will be performed sequentially, thus races are avoided.

----- RATIONALE -----

>> A5: Why did you choose the data structure(s) that you did for  
>> representing virtual-to-physical mappings?  
We added "upage" and "frame" in spte\_entry and frame\_entry  
correspondingly.  
As described in A3, only user virtual address is used and kernel  
virtual address  
is never used, thus alias problem can be avoided.

PAGING TO AND FROM DISK  
=====

----- DATA STRUCTURES -----

>> B1: Copy here the declaration of each new or changed `struct'  
or  
>> `struct' member, global or static variable, `typedef', or  
>> enumeration. Identify the purpose of each in 25 words or  
less.

In "swap.h":  
struct block \*swap\_block; // a device where swapped pages are  
located  
struct bitmap \*swap\_bitmap; // indicate whether the sectors on  
swap block is occupied or free  
struct lock swap\_lock;

----- ALGORITHMS -----

>> B2: When a frame is required but none is free, some frame  
must be  
>> evicted. Describe your code for choosing a frame to evict.  
We used a clock algorithm to choose which frame to evict.  
Going through all the frames in frame table:  
1. If the page is not pinned, go to Step 2, otherwise,  
continue to check the next frame;  
2. Check the frame, if the page has been accessed recently,  
set the accessed bit to false and continue to check the next  
frame(Go to step 1);  
3. Otherwise, if the page has a dirty bit, there are two  
cases:  
1> if it is a mmap file, write it back to the file  
2> if it is executable file, swap it out.

Go to step 4;

4. Remove this frame from the frame list, free the page and set loaded to false;

5. Call `palloc_get_page(flag)` and return a free frame.

>> B3: When a process P obtains a frame that was previously used by a

>> process Q, how do you adjust the page table (and any other data

>> structures) to reflect the frame Q no longer has?

When a frame is freed or evicted, we will call

`pagedir_clear_page()` and `palloc_free_page()` to clear it.

Then if the frame is obtained by other processes, we will

set the "owner" in frame entry to `current_thread()`,

set new "spte" and new frame address. In this way, no old process information will be left.

>> B4: Explain your heuristic for deciding whether a page fault for an

>> invalid virtual address should cause the stack to be extended into

>> the page that faulted.

If an supplemental page is not found in page table, we will call `grow_stack()`.

We set the heuristic to be 32, because `PUSHA` might cause an access to be 32 bytes below stack pointer. So we check the user memory address, if it is between `PHYS_BASE` and `esp - 32`, it is valid.

But the max size of a stack is 8MB, so it is beyond this size, it is invalid.

(refer to <https://github.com/ryantimwilson/Pintos-Project-3/blob/dd973a541d9c6b4009b00423e05a1f377ef5c11a/src/userprog/syscall.h#L14>)

---- SYNCHRONIZATION ----

>> B6: A page fault in process P can cause another process Q's frame

>> to be evicted. How do you ensure that Q cannot access or

modify

>> the page during the eviction process? How do you avoid a race

>> between P evicting Q's frame and Q faulting the page back in?

As we mentioned in B2, a process will only evict a frame when its

"pinned" field is false.

So when process Q is using frame A, it will set its "pinned" field

to true, thus process P cannot evict this frame.

If Q is not using frame A, then process P can evict it.

After evicting it, process P will set the "loaded" field to be false, so next time, when process Q wants to use the supplemental

page, it needs to reload it.

>> B7: Suppose a page fault in process P causes a page to be read from

>> the file system or swap. How do you ensure that a second process Q

>> cannot interfere by e.g. attempting to evict the frame while it is

>> still being read in?

If a frame is still being read in, its "pinned" field will be set to true, so process P cannot evict it.

>> B8: Explain how you handle access to paged-out pages that occur

>> during system calls. Do you use page faults to bring in pages (as

>> in user programs), or do you have a mechanism for "locking" frames

>> into physical memory, or do you use some other design? How do you

>> gracefully handle attempted accesses to invalid virtual addresses?

During system calls, we will check the pointer before we call any function. So if there exists a page fault, we will call `load_page()` or `grow_stack()` to get the page. We also set "pinned" to lock the pages. If a virtual address is invalid, process will exit.

#### ---- RATIONALE ----

>> B9: A single lock for the whole VM system would make  
>> synchronization easy, but limit parallelism. On the other  
hand,  
>> using many locks complicates synchronization and raises the  
>> possibility for deadlock but allows for high parallelism.  
Explain  
>> where your design falls along this continuum and why you  
chose to  
>> design it this way.

We use several locks in order to achieve parallelism. As  
described  
before, we have a boolean field to detect some page faults.

#### MEMORY MAPPED FILES =====

#### ---- DATA STRUCTURES ----

>> C1: Copy here the declaration of each new or changed `struct'  
or  
>> `struct' member, global or static variable, `typedef', or  
>> enumeration. Identify the purpose of each in 25 words or  
less.

In "page.h":

```
// mmap_entry is an entry to record the supplemental page of
// a memory mapped file.
struct mmap_entry {
    struct spt_entry *spte;    // point out the related
    supplemental page table
    int mapid;                // map id
    struct list_elem elem;
};
```

In "exception.c":

```
struct list mmap_list; // list of mmap entry;
int mapid;              // Incremented every time when a file is
mapped.
```

#### ---- ALGORITHMS ----

>> C2: Describe how memory mapped files integrate into your  
virtual

>> memory subsystem. Explain how the page fault and eviction  
>> processes differ between swap pages and other pages.

Memory mapped files are lazy loaded. When a page fault occurs, it will read from the file just like the executable files. This is the same for loading page. The only difference is in eviction and munmap.

Mmap files are recorded in the mmap\_list, so in munmap, we get the page from mmap\_list, and write dirty page back to disk.

>> C3: Explain how you determine whether a new file mapping overlaps  
>> any existing segment.

The file is mapped page by page. When a new page is mapped, before adding it to the page table, we will check if it already exists or not. If it already exists, return false. Otherwise, create a new map file page, and call hash\_insert() to insert it to the page table.

---- RATIONALE ----

>> C4: Mappings created with "mmap" have similar semantics to those of  
>> data demand-paged from executables, except that "mmap" mappings are  
>> written back to their original files, not to swap. This implies  
>> that much of their implementation can be shared. Explain why your  
>> implementation either does or does not share much of the code for  
>> the two situations.

Both mmap file and executable files are lazy load, so we use the same load function in these two situations. But in frame\_evict(), if the page has a dirty bit and is a mmap file, we write it back to disk. Otherwise, if it is an executable file, we call swap\_out(), this is an extension.

SURVEY QUESTIONS  
=====



Answering these questions is optional, but it will help us improve the course in future quarters. Feel free to tell us anything you want--these questions are just to spur your thoughts. You may also choose to respond anonymously in the course evaluations at the end of the quarter.

>> In your opinion, was this assignment, or any one of the three problems

>> in it, too easy or too hard? Did it take too long or too little time?

>> Did you find that working on a particular part of the assignment gave

>> you greater insight into some aspect of OS design?

>> Is there some particular fact or hint we should give students in

>> future quarters to help them solve the problems? Conversely, did you

>> find any of our guidance to be misleading?

>> Do you have any suggestions for the TAs to more effectively assist

>> students, either for future quarters or the remaining projects?

>> Any other comments?